

Objective(s):

This activity aims to perform regression analysis using polynomial regression

Intended Learning Outcomes (ILOs):

- Demonstrate how to build a regression model to predict the outcome using polynomial regression.
- Evaluate the performance of the regression model using polynomial regression

Resources:

- Jupyter Notebook
- internet_traffic_hist.csv

Procedure:

Using numpy polyfit to perform polynomial regression

Import the libraries and the data

```
In [127... from google.colab import drive  
  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [128... import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [129... from sklearn.metrics import r2_score  
from scipy.optimize import curve_fit
```

Load the dataset

```
In [130... internet = '/content/drive/MyDrive/DATASETS/internet_traffic_hist-2.csv'  
  
df_hist = pd.read_csv(internet)  
df_hist.head(11)
```

```
Out[130]:
```

	traffic	year
0	100.000000	2005
1	126.933755	2006
2	160.303757	2007
3	203.390603	2008
4	241.292566	2009
5	308.791823	2010
6	379.980659	2011
7	495.840568	2012
8	616.207252	2013
9	752.103483	2014
10	931.200929	2015

Build the first order polynomial using numpy polyfit

In [131...

```
order = 1

# XY Plot of year and traffic
x = df_hist.year
y = df_hist.traffic

m, b = np.polyfit(x,y,order)

plt.plot(x, y, label = 'Historical Internet Traffic', linewidth = 7)
plt.plot(x, y, '*k', markersize = 15, label = '')
plt.plot(x, m*x + b, '-', label = 'Simple Linear Regression Line', linewidth = 6)

print ('The slope of line is {}'.format(m))
print ('The y intercept is {}'.format(b))
print ('The best fit simple linear regression line is {}x + {}'.format(m,b))

#Increase slightly the axis sizes to make the plot more clear
plt.axis([x.iloc[0]-1, x.iloc[-1]+1, y.iloc[0]*-0.1, y.iloc[-1]*1.1])

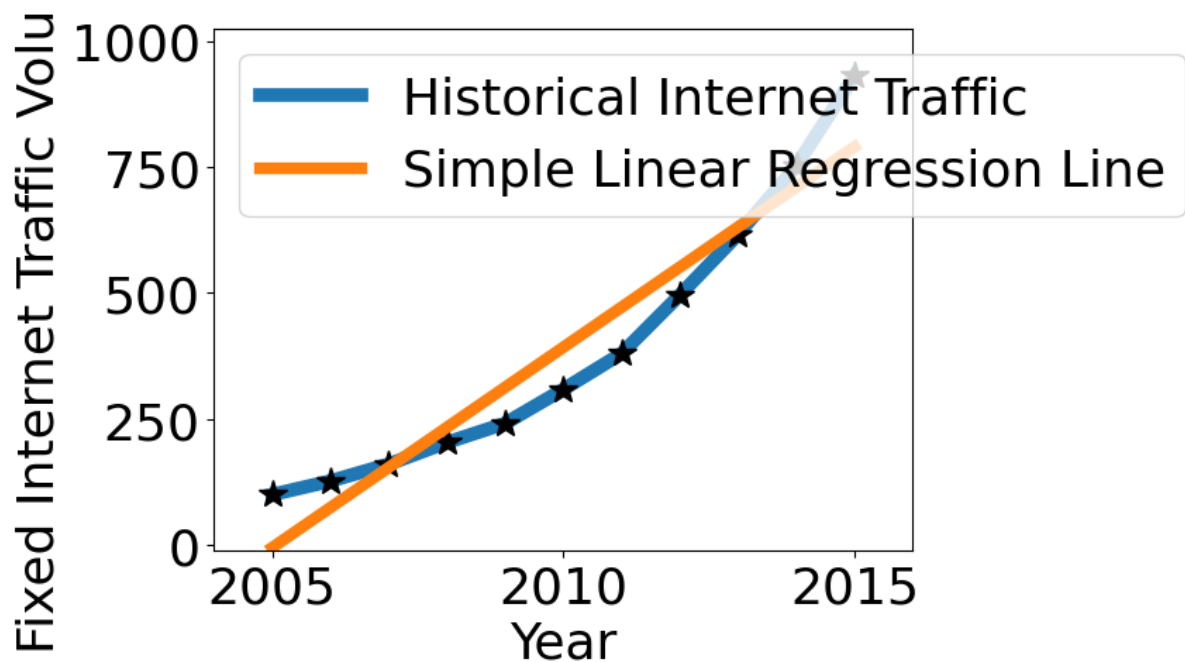
# Add axis labels
plt.xlabel('Year')
plt.ylabel('Fixed Internet Traffic Volume')
plt.legend(loc = 'upper left')

# Increase default font size
plt.rcParams.update({'font.size': 26})
plt.show()
```

The slope of line is 79.52710966244513.

The y intercept is -159457.12265833947.

The best fit simple linear regression line is 79.52710966244513x + -159457.12265833947.



Build the model using Higher Order Polynomial (1 to 4)

In [132...

```
models = []          # to store polynomial model parameters (list of poly1d objects)
errors_hist = []     # to store the absolute errors for each point (2005-2015) and for
mse_hist = []        # to store the MSE for each model (list of numpy floats)

# Try polynomial models with increasing order
for order in range(1,4):
    # Fit polynomial model
    p = (np.poly1d(np.polyfit(x, y, order)))
    models.append(p)

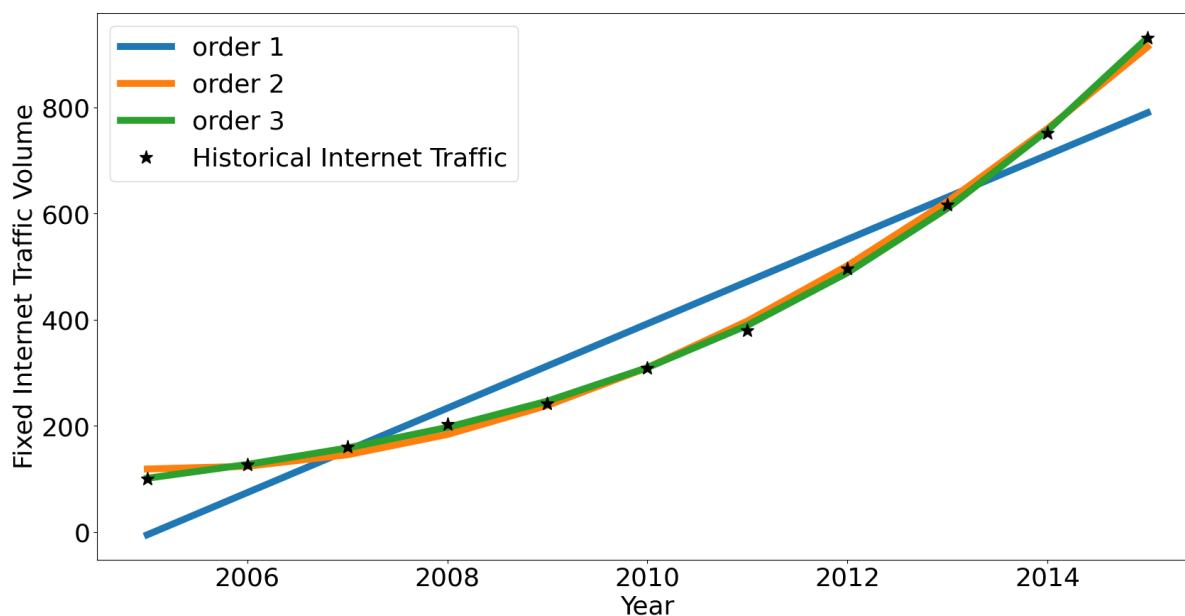
plt.figure(figsize = (20,10))

# Visualize polynomial models fit
for model in models[0:3]:
    plt.plot(x, model(x), label = 'order ' + str(len(model)), linewidth = 7)

plt.plot(x, y, '*k', markersize = 14, label = 'Historical Internet Traffic', linewidth = 2)
plt.legend(loc = 'upper left')

# Add axis labels
plt.xlabel('Year')
plt.ylabel('Fixed Internet Traffic Volume')

plt.show()
```



Calculate the error for each order

In [133...

```
# Calculate and store the errors
models = []          # to store polynomial model parameters (list of poly1d objects)
errors_hist = []     # to store the absolute errors for each point (2005-2015) and for
mse_hist = []        # to store the MSE for each model (list of numpy floats)

#Try polynomial models with increasing order
for order in range(1,4):
    # Fit polynomial model
    p = (np.poly1d(np.polyfit(x, y, order)))
    models.append(p)

    e = np.abs(y-p(x))          # absolute error
    mse = np.sum(e**2)/len(df_hist) # mse

    errors_hist.append(e)      #Store the absolute errors
    mse_hist.append(mse)      # Store the mse
```

In [134...

```

# Visualize fit error for each year

x = df_hist.year
width = 0.2 #size of the bar

fig = plt.figure(figsize=(20,10))
ax = fig.add_subplot(111)

p1 = ax.bar( x, errors_hist[0], width, color = 'b', label = 'Abs. error order 1 fit
p2 = ax.bar( x + width, errors_hist[1], width, color = 'r', label = 'Abs. error ord
p3 = ax.bar( x + 2*width, errors_hist[2], width, color = 'y', label = 'Abs. error o

# "Prettyfy" the bar graph
ax.set_xticks(x+2*width)
ax.set_xticklabels(x)
plt.legend(loc = 'upper left', fontsize =16)
plt.show()

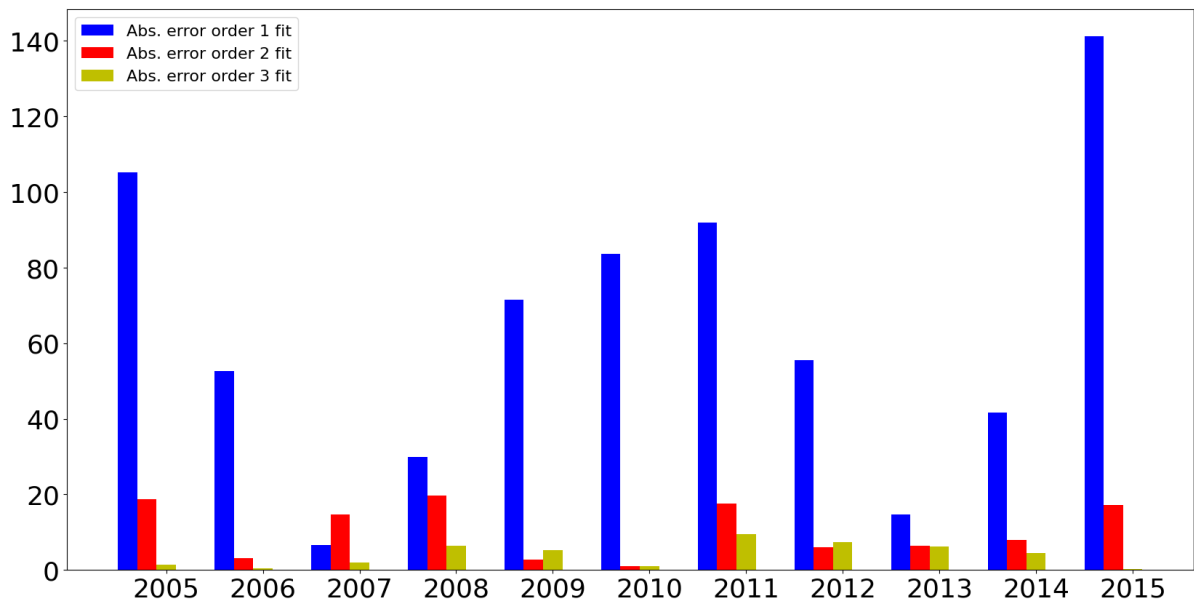
#Visualise MSE for each model
fig = plt.figure(figsize=(20,10))
ax = fig.add_subplot(111)

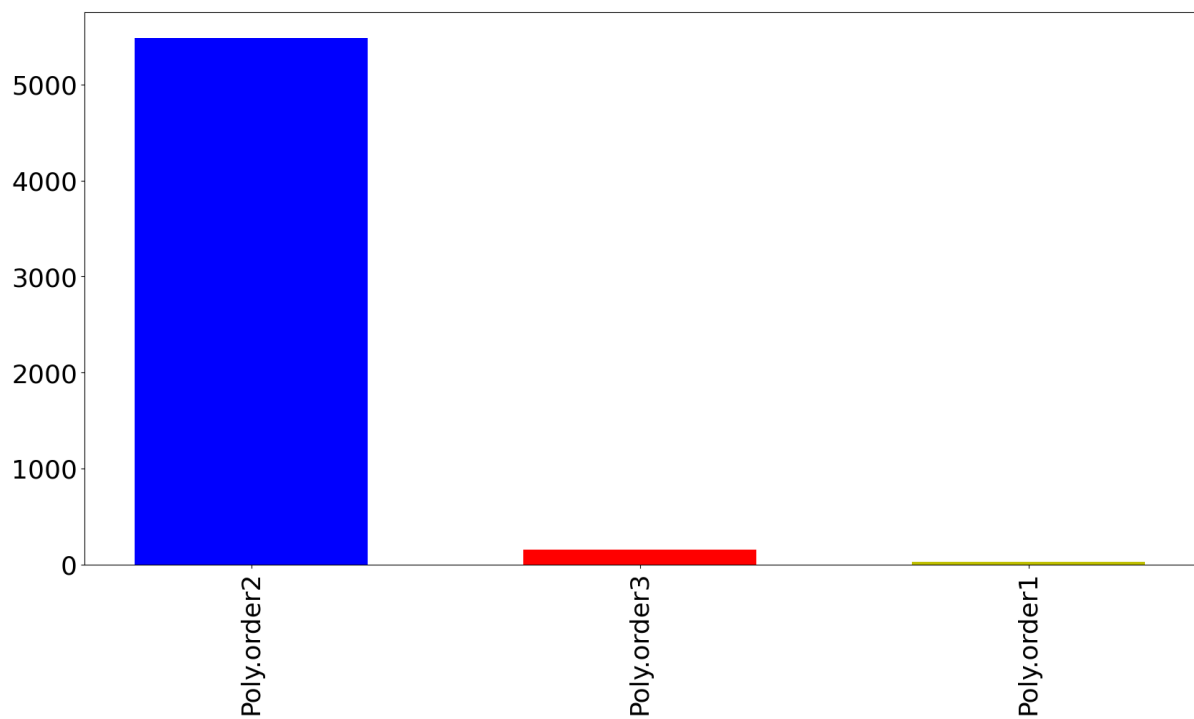
x = np.array([0,1,2,3])
width = .6 #size of the bar

p1 = ax.bar( x[0], mse_hist[0], width, color = 'b', label = 'pred. error order 1 fi
p2 = ax.bar( x[1], mse_hist[1], width, color = 'r', label = 'pred. error order 2 fi
p3 = ax.bar( x[2], mse_hist[2], width, color = 'y', label = 'pred. error order 3 fi

ax.set_xticks(ticks=[0.0,1.0,2.0],labels={'Poly.order1', 'Poly.order2', 'Poly.order
plt.show()

```





Interpret the result of the fit error for each year

my interpretation of the fit error per order in the our polynomial regression is that as the **order of the polynomial regression (from the code above) increases** the **error between the model and the true value decreases**

In [135...

```

# Polynomial function order
order = 3

x = df_hist.year.values      # regressor
y = df_hist.traffic.values   # regressand

# Fit the model, return the polynomial parameter values in a numpy array such that
# y = p[0]*x**order + p[1]*x*(order-1) ...

p_array = np.polyfit(x,y,order)

print(type(p_array), p_array)

# poly1d is a convenience class, used to encapsulate "natural" operations on polyno
# so that said operations may take on their customary form in code

# wrap the p_array in a poly1 object
p = np.poly1d(p_array)

print(type(p), p)

# use the poly1d object to evaluate the value of the polynomial in a specific point
print('The value of the polynomial for x = 2020 is : {}'.format(p(2020)))

# compute the absolute error for each value of x and the MSE error for the estimate
e = np.abs(y-p(x))
mse = np.sum(e**2)/len(x)

print('The estimated polynomial parameters are: {}'.format(p))
print('The errors for each value of x, given the estimated polynomial parameters ar
print('The MSE is :{}'.format(mse))

<class 'numpy.ndarray'> [ 4.83129404e-01 -2.90500578e+03  5.82252085e+06 -3.8900538
7e+09]
<class 'numpy.poly1d'>      3      2
0.4831 x - 2905 x + 5.823e+06 x - 3.89e+09
The value of the polynomial for x = 2020 is : 2328.5784521102905
The estimated polynomial parameters are:      3      2
0.4831 x - 2905 x + 5.823e+06 x - 3.89e+09
The errors for each value of x, given the estimated polynomial parameters are:
[1.30743027 0.39125264 2.02722693 6.32983208 5.28394403 0.93069802
 9.41692212 7.34010081 6.27729748 4.48133933 0.16291521]
The MSE is :25.17218620372407

```

Using sklearn to perform polynomial regression

Import the necessary libraries

In [136...

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

```

Training the Polynomial Regression model using degree 3

In [137...

```

poly_reg = PolynomialFeatures(degree=3)
X_poly = poly_reg.fit_transform(x.reshape(-1, 1))

```

```
In [138... lin_reg = LinearRegression()  
lin_reg.fit(X_poly,y)
```

```
Out[138]: ▼ LinearRegression  
LinearRegression()
```

Predict the result using polynomial regression model

```
In [139... y_pred = lin_reg.predict(X_poly)
```

```
In [140... df = pd.DataFrame({'Real Values': y, 'Predicted Values':y_pred})
```

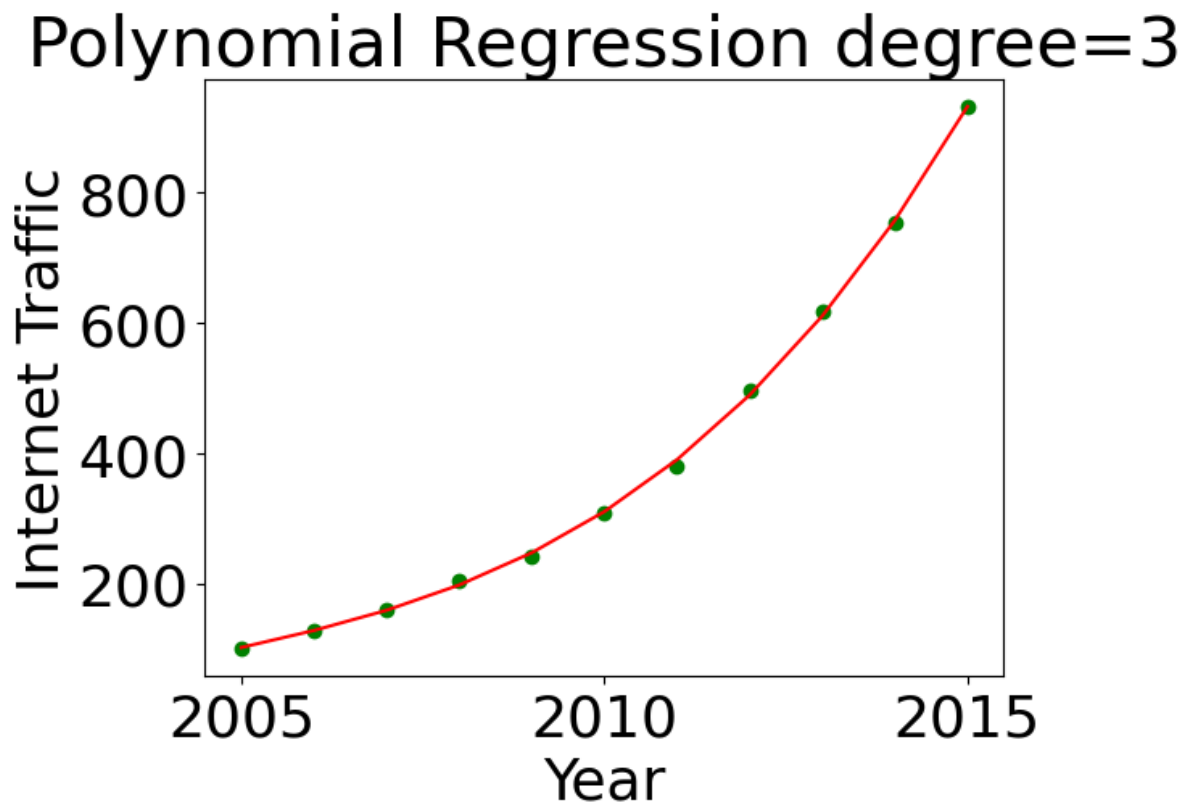
```
In [141... df
```

```
Out[141]:
```

	Real Values	Predicted Values
0	100.000000	101.241620
1	126.933755	127.295559
2	160.303757	158.270686
3	203.390603	197.068289
4	241.292566	246.589652
5	308.791823	309.736069
6	379.980659	389.408823
7	495.840568	488.509205
8	616.207252	609.938500
9	752.103483	756.598000
10	931.200929	931.388989

Visualize the Polynomial Regression results

```
In [142... plt.scatter(x, y, color='green')  
plt.plot(x, y_pred, color = 'red')  
  
plt.title("Polynomial Regression degree=3")  
plt.xlabel('Year')  
plt.ylabel('Internet Traffic')  
  
plt.show();
```

Supplementary Activity:

- Choose your own dataset
- Import the dataset
- Perform polynomial regression using sklearn and polyfit
- Measure the performance for each polynomial degree.
- Plot the performance of the model for each polynomial degree.

In [143]...

```
suppath = '/content/drive/MyDrive/DATASETS/Advertising Budget and Sales.csv'
suppdf = pd.read_csv(suppath)
suppdf.head()
```

Out[143]:

	Unnamed: 0	TV Ad Budget (\$)	Radio Ad Budget (\$)	Newspaper Ad Budget (\$)	Sales (\$)
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

In [144]...

```
# cleaning the data
suppdf.drop(columns=['Unnamed: 0'], inplace=True)
suppdf.head()
```

Out[144]:

	TV Ad Budget (\$)	Radio Ad Budget (\$)	Newspaper Ad Budget (\$)	Sales (\$)
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

In [145]...

```
# renaming columns
suppdf.rename(columns={'TV Ad Budget ($)':'TVBudget(USD)',
                       'Radio Ad Budget ($)':'RadioBudget(USD)',
                       'Newspaper Ad Budget ($)':'NewspaperBudget(USD)',
                       'Sales ($)':'Sales(USD)'}, inplace=True)

suppdf.head()
```

Out[145]:

	TVBudget(USD)	RadioBudget(USD)	NewspaperBudget(USD)	Sales(USD)
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

In [146]...

```
# checking for null values
suppdf.isnull().sum()
```

Out[146]:

	0
TVBudget(USD)	0
RadioBudget(USD)	0
NewspaperBudget(USD)	0
Sales(USD)	0

dtype: int64

In [147]...

```
# getting the independent variable and dependent variable
x = suppdf['TVBudget(USD)'].values/np.max(suppdf['TVBudget(USD)'].values)
y = suppdf['Sales(USD)']
```

```

In [148... #fitting and graphing the model
models = [] # to store polynomial model parameters (list of poly1d objects)

for order in range(1,4):# loop for fitting the polynomial regression per degree
    p = (np.poly1d(np.polyfit(x, y, order)))
    models.append(p)

plt.figure(figsize = (20,10))

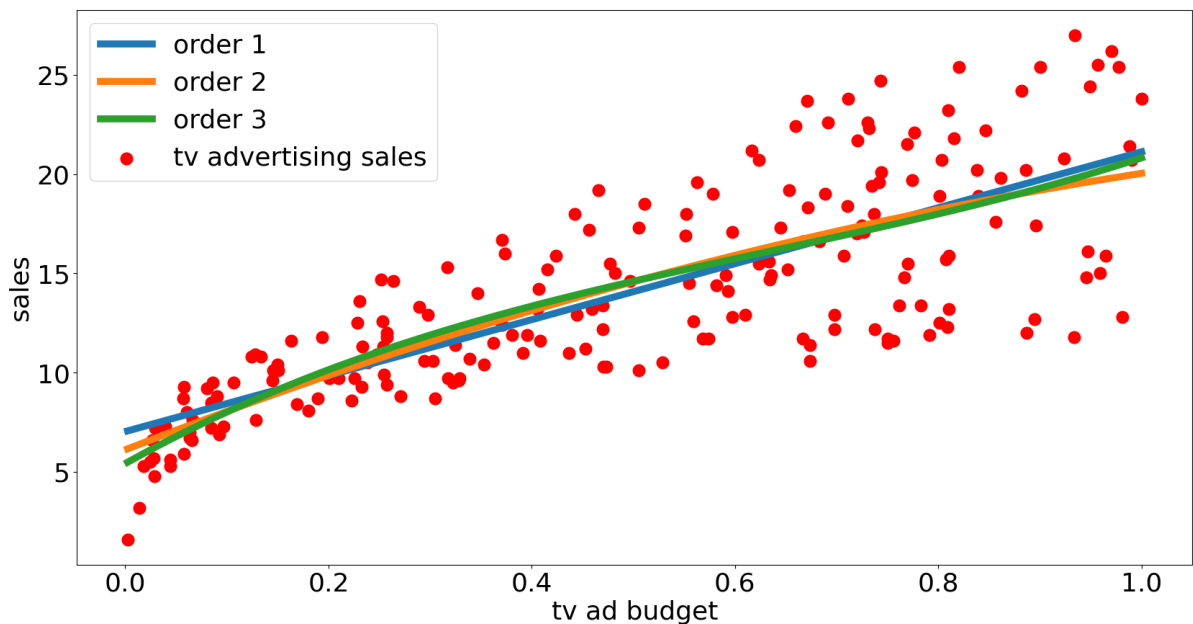
for model in models[0:3]: # loop for graphing the model per order in polynomial reg
    plt.plot(sorted(x), model(x)[np.argsort(x.ravel())], label = 'order ' + str(len

plt.scatter(x, y, label = 'tv advertising sales', linewidth = 7,color='red')
plt.legend(loc = 'upper left')

# Adding axis Labels
plt.xlabel('tv ad budget')
plt.ylabel('sales')

```

Out[148]: Text(0, 0.5, 'sales')



```

In [149... mse_hist = [] # to store the MSE for each model (list of numpy floats)

#Try polynomial models with increasing order
for model in models[0:3]:
    mse = np.sum(e**2)/len(df_hist) # get the mean error of every order model
    mse_hist.append(mse) # store the mse in the list

```

```

In [166... ord = 1
for order in range(0,3):
    print(f'order {ord} mean error:', mse_hist[order])
    ord += 1

```

```

order 1 mean error: 191.13914392103197
order 2 mean error: 187.61205509549927
order 3 mean error: 186.15210094717514

```

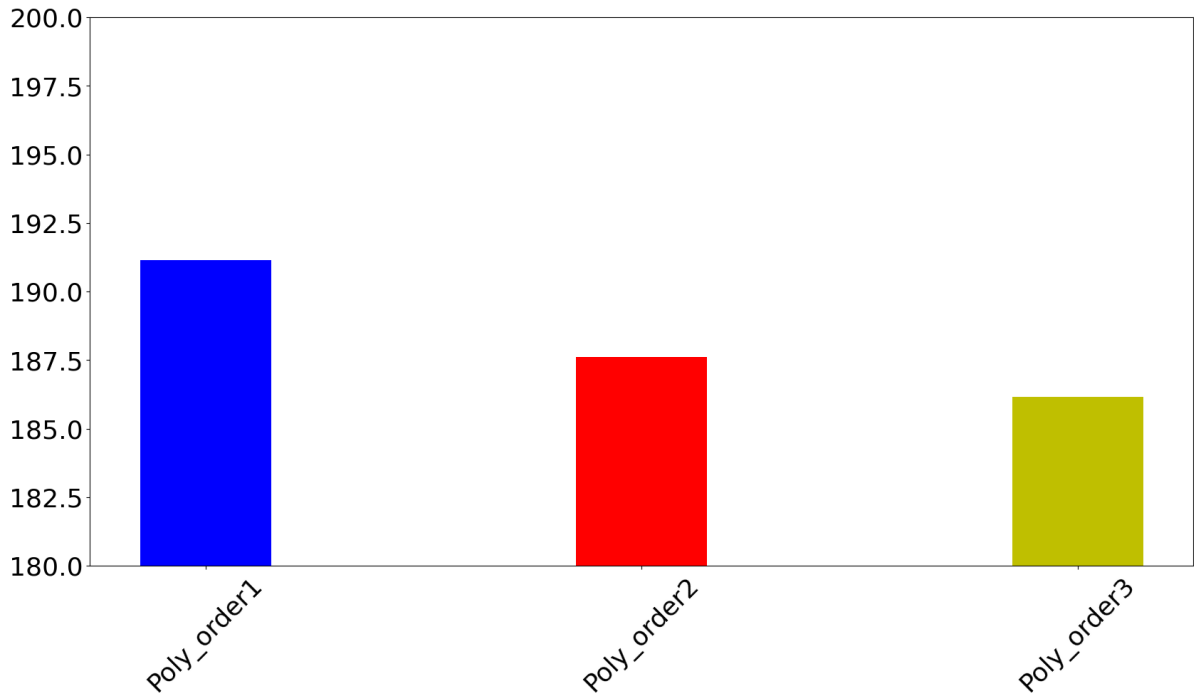
In [168...]

```
#graphing the mean error of the model we created
fig = plt.figure(figsize=(20,10))
ax = fig.add_subplot(111)

x = np.array([0,1,2])
p1 = ax.bar( x[0], mse_hist[0], width, color = 'b', label = 'pred. error order 1 fi
p2 = ax.bar( x[1], mse_hist[1], width, color = 'r', label = 'pred. error order 2 fi
p3 = ax.bar( x[2], mse_hist[2], width, color = 'y', label = 'pred. error order 3 fi

plt.ylim(180,200) #zooms in the value of y limit
ax.set_xticks(ticks=[0.0,1.0,2.0],labels=['Poly_order1', 'Poly_order2', 'Poly_order
```

```
Out[168]: [<matplotlib.axis.XTick at 0x7b6fbbc8edd0>,
<matplotlib.axis.XTick at 0x7b6fbbc8d9f0>,
<matplotlib.axis.XTick at 0x7b6fbbc8e560>]
```



Conclusion:

what I learn today in this activity is about polynomial regression, how to model it, evaluate it by getting the mean error, and the absolute error, I learned that by fitting the polynomial comes with the value called order the value of the order of the poly regression is inversely proportional to the error of the model, I also concluded that as order of the poly fit increases the more it joins the trend of the datapoints of our dataset